

ASSIGNMENT-4 (Anish Singh) 21BCE3775

Logistic regression, Decision tree and random forest classifiers on Employee Attrition dataset

Data Preprocessing.

```
In [1]: #Importing necessary libraries.
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
```


```
In [2]: #Importing the dataset.
df=pd.read_csv("Employee-Attrition.csv")
```

```
In [3]: df.head()
```

```
Out[3]:
```

	Age	Attrition	BusinessTravel	DailyRate	Department	DistanceFromHome	Education	Educ
0	41	Yes	Travel_Rarely	1102	Sales	1	2	Lif
1	49	No	Travel_Frequently	279	Research & Development	8	1	Lif
2	37	Yes	Travel_Rarely	1373	Research & Development	2	2	
3	33	No	Travel_Frequently	1392	Research & Development	3	4	Lif
4	27	No	Travel_Rarely	591	Research & Development	2	1	

5 rows × 35 columns



```
In [4]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 1470 entries, 0 to 1469
Data columns (total 35 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   Age                                   1470 non-null   int64
1   Attrition                           1470 non-null   object
2   BusinessTravel                       1470 non-null   object
3   DailyRate                           1470 non-null   int64
4   Department                           1470 non-null   object
5   DistanceFromHome                    1470 non-null   int64
6   Education                           1470 non-null   int64
7   EducationField                       1470 non-null   object
8   EmployeeCount                       1470 non-null   int64
9   EmployeeNumber                      1470 non-null   int64
10  EnvironmentSatisfaction              1470 non-null   int64
11  Gender                              1470 non-null   object
12  HourlyRate                          1470 non-null   int64
13  JobInvolvement                      1470 non-null   int64
14  JobLevel                            1470 non-null   int64
15  JobRole                             1470 non-null   object
16  JobSatisfaction                     1470 non-null   int64
17  MaritalStatus                       1470 non-null   object
18  MonthlyIncome                       1470 non-null   int64
19  MonthlyRate                         1470 non-null   int64
20  NumCompaniesWorked                  1470 non-null   int64
21  Over18                             1470 non-null   object
22  OverTime                            1470 non-null   object
23  PercentSalaryHike                   1470 non-null   int64
24  PerformanceRating                   1470 non-null   int64
25  RelationshipSatisfaction             1470 non-null   int64
26  StandardHours                       1470 non-null   int64
27  StockOptionLevel                    1470 non-null   int64
28  TotalWorkingYears                   1470 non-null   int64
29  TrainingTimesLastYear               1470 non-null   int64
30  WorkLifeBalance                     1470 non-null   int64
31  YearsAtCompany                      1470 non-null   int64
32  YearsInCurrentRole                  1470 non-null   int64
33  YearsSinceLastPromotion              1470 non-null   int64
34  YearsWithCurrManager                1470 non-null   int64
dtypes: int64(26), object(9)
memory usage: 402.1+ KB
```

```
In [5]: #Checking for Null Values.  
df.isnull().any()
```

```
Out[5]: Age                False  
Attrition                 False  
BusinessTravel            False  
DailyRate                 False  
Department                False  
DistanceFromHome          False  
Education                 False  
EducationField             False  
EmployeeCount              False  
EmployeeNumber             False  
EnvironmentSatisfaction    False  
Gender                     False  
HourlyRate                 False  
JobInvolvement             False  
JobLevel                   False  
JobRole                     False  
JobSatisfaction            False  
MaritalStatus              False  
MonthlyIncome              False  
MonthlyRate                False  
NumCompaniesWorked         False  
Over18                     False  
OverTime                   False  
PercentSalaryHike          False  
PerformanceRating          False  
RelationshipSatisfaction    False  
StandardHours              False  
StockOptionLevel           False  
TotalWorkingYears          False  
TrainingTimesLastYear      False  
WorkLifeBalance            False  
YearsAtCompany              False  
YearsInCurrentRole          False  
YearsSinceLastPromotion     False  
YearsWithCurrManager        False  
dtype: bool
```

```
In [6]: df.isnull().sum()
```

```
Out[6]: Age                                0
Attrition                                0
BusinessTravel                           0
DailyRate                                0
Department                               0
DistanceFromHome                         0
Education                                0
EducationField                           0
EmployeeCount                            0
EmployeeNumber                           0
EnvironmentSatisfaction                  0
Gender                                   0
HourlyRate                               0
JobInvolvement                           0
JobLevel                                 0
JobRole                                  0
JobSatisfaction                          0
MaritalStatus                            0
MonthlyIncome                           0
MonthlyRate                              0
NumCompaniesWorked                       0
Over18                                   0
OverTime                                  0
PercentSalaryHike                        0
PerformanceRating                        0
RelationshipSatisfaction                  0
StandardHours                            0
StockOptionLevel                         0
TotalWorkingYears                       0
TrainingTimesLastYear                    0
WorkLifeBalance                          0
YearsAtCompany                           0
YearsInCurrentRole                       0
YearsSinceLastPromotion                   0
YearsWithCurrManager                     0
dtype: int64
```

```
In [7]: #Data Visualization.  
sns.distplot(df["Age"])
```

C:\Users\Admin\AppData\Local\Temp\ipykernel_39480\2400079689.py:2: UserWarning:

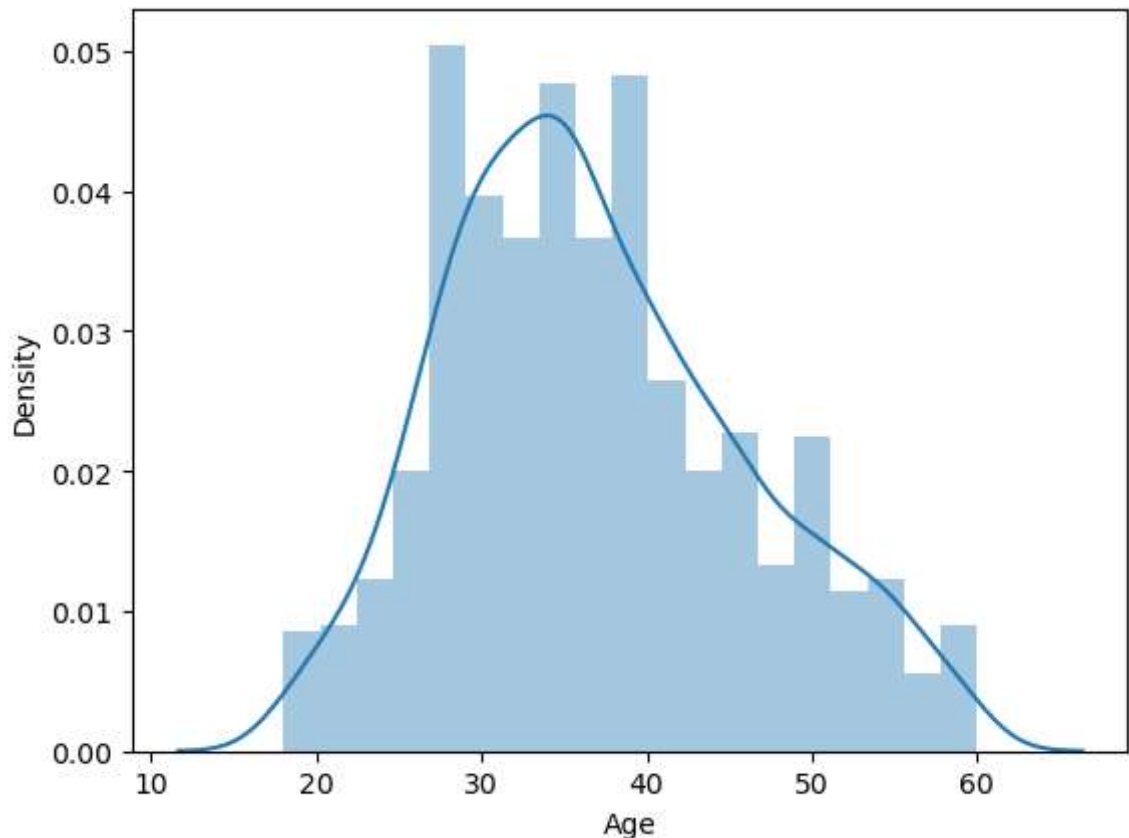
`distplot` is a deprecated function and will be removed in seaborn v0.14.0.

Please adapt your code to use either `displot` (a figure-level function with similar flexibility) or `histplot` (an axes-level function for histograms).

For a guide to updating your code to use the new functions, please see <https://gist.github.com/mwaskom/de44147ed2974457ad6372750bbe5751> (<https://gist.github.com/mwaskom/de44147ed2974457ad6372750bbe5751>)

```
sns.distplot(df["Age"])
```

Out[7]: <Axes: xlabel='Age', ylabel='Density'>



```
In [8]: attrition_count = pd.DataFrame(df['Attrition'].value_counts())  
plt.pie(attrition_count['Attrition'], labels = ['No', 'Yes'], explode = (0.2, 0.2))
```

```
Out[8]: ([<matplotlib.patches.Wedge at 0x2634cd0cb80>,  
<matplotlib.patches.Wedge at 0x2634ce105e0>],  
[Text(-1.136781068348268, 0.6306574368426737, 'No'),  
Text(0.961891673217765, -0.5336332157899547, 'Yes')])
```




```
In [9]: plt.figure(figsize=[20,20])
sns.heatmap(df.corr(),annot=True)
```

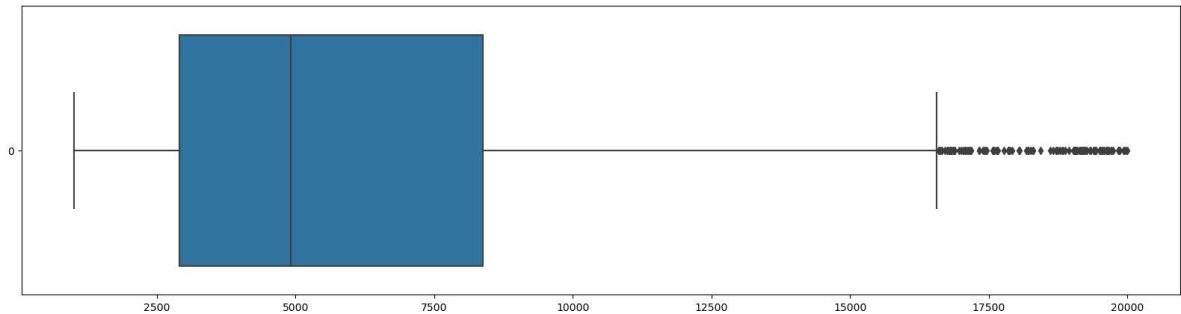
C:\Users\Admin\AppData\Local\Temp\ipykernel_39480\3113117044.py:2: FutureWarning: The default value of numeric_only in DataFrame.corr is deprecated. In a future version, it will default to False. Select only valid columns or specify the value of numeric_only to silence this warning.

```
sns.heatmap(df.corr(),annot=True)
```


Out[9]: <Axes: >



```
In [10]: #Outlier detection
plt.figure(figsize=[20,5])
sns.boxplot(df['MonthlyIncome'],orient='h')
plt.show()
```



```
In [11]: plt.figure(figsize=[20,5])
sns.boxplot(df['YearsAtCompany'],orient='h')
plt.show()
```



```
In [12]: # Label Encoding
categories = ['BusinessTravel','Department','Education','EducationField','Gender','EnvironmentSatisfaction','JobInvolvement','JobLevel','JobRole','PerformanceRating','RelationshipSatisfaction','StockOptionLevel']
categorical = df[categories].astype('object')
categorical = pd.get_dummies(df[categories], drop_first = True)
```

```
In [13]: # Splitting Dependent and Independent variables
independent = ['Attrition','Over18','EmployeeCount','StandardHours','EmployeeNumber']
continuous = df.drop(columns= categories)
continuous = continuous.drop(columns= independent)
```

```
In [14]: # X - Features, Y- Target variables
X = pd.concat([categorical,continuous],axis=1)
Y = df['Attrition'].replace({'Yes': 1, 'No': 0}).values.reshape(-1,1)
```



```
In [15]: # Feature scaling
from sklearn.preprocessing import StandardScaler

scaler = StandardScaler()

continuous_variables = list(continuous.columns)

X = X.reset_index()
del X['index']
X[continuous_variables] = pd.DataFrame(scaler.fit_transform(X[continuous_varia
```


```
In [16]: #Splitting Data into Train and Test.
from sklearn.model_selection import train_test_split
x_train,x_test,y_train,y_test=train_test_split(X,Y,test_size=0.2,random_state=
```

```
In [17]: x_train.shape,x_test.shape,y_train.shape,y_test.shape
```

```
Out[17]: ((1176, 44), (294, 44), (1176, 1), (294, 1))
```

Logistic Regression model

```
In [18]: #Importing necessary libraries
from sklearn.linear_model import LogisticRegression
from sklearn.metrics import accuracy_score,precision_score, recall_score, f1_s
```



```
In [19]: #Initializing the model
lr = LogisticRegression()
```

```
In [20]: #Training the model
lr.fit(x_train,y_train)
```

C:\Users\Admin\anaconda3\lib\site-packages\sklearn\utils\validation.py:1143:
DataConversionWarning: A column-vector y was passed when a 1d array was expected. Please change the shape of y to (n_samples,), for example using ravel().

y = column_or_1d(y, warn=True)
C:\Users\Admin\anaconda3\lib\site-packages\sklearn\linear_model_logistic.py:
458: ConvergenceWarning: lbfgs failed to converge (status=1):
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.

Increase the number of iterations (max_iter) or scale the data as shown in:
<https://scikit-learn.org/stable/modules/preprocessing.html> (<https://scikit-learn.org/stable/modules/preprocessing.html>)
Please also refer to the documentation for alternative solver options:
https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression (https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression)
n_iter_i = _check_optimize_result(

```
Out[20]: LogisticRegression()
```

In a Jupyter environment, please rerun this cell to show the HTML representation or trust the notebook.

On GitHub, the HTML representation is unable to render, please try loading this page with nbviewer.org.

```
In [21]: #Testing the model
y_pred = lr.predict(x_test)
```

```
In [22]: # Evaluation of model
# Accuracy score
print("Accuracy of Logistic regression model:",accuracy_score(y_test,y_pred))
```

Accuracy of Logistic regression model: 0.8843537414965986

```
In [23]: # Precision score
precision_yes = precision_score(y_test, y_pred, pos_label=1)
print("Precision (Yes): " + str(round(precision_yes, 2)))
precision_no = precision_score(y_test, y_pred, pos_label=0)
print("Precision (No): " + str(round(precision_no, 2)))
```

Precision (Yes): 0.76
Precision (No): 0.9

```
In [24]: # Recall score
recall_yes = recall_score(y_test, y_pred, pos_label=1)
print("Recall (Yes): " + str(round(recall_yes, 2)))
recall_no = recall_score(y_test, y_pred, pos_label=0)
print("Recall (No): " + str(round(recall_no, 2)))
```

Recall (Yes): 0.45
Recall (No): 0.97

```
In [25]: # F1 score
f1_score_yes = f1_score(y_test, y_pred, pos_label=1)
print("F1 Score (Yes): " + str(round(f1_score_yes, 2)))
f1_score_no = f1_score(y_test, y_pred, pos_label=0)
print("F1 Score (No): " + str(round(f1_score_no, 2)))
```

F1 Score (Yes): 0.56

F1 Score (No): 0.93

```
In [26]: # Confusion matrix
print("Confusion matrix:\n\n",confusion_matrix(y_test,y_pred))
```

Confusion matrix:


```
[[238   7]
 [ 27  22]]
```

```
In [27]: # Classification Report
print("Classification report of Logistic Regression model:\n\n",classification_report(y_test,y_pred))
```

Classification report of Logistic Regression model:

	precision	recall	f1-score	support
0	0.90	0.97	0.93	245
1	0.76	0.45	0.56	49
accuracy			0.88	294
macro avg	0.83	0.71	0.75	294
weighted avg	0.87	0.88	0.87	294

```
In [28]: # ROC curve
probability = lr.predict_proba(x_test)[:,-1]
fpr, tpr, thresholds = roc_curve(y_test, probability)
plt.plot(fpr, tpr)
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
plt.title('ROC Curve of Logistic regression model')
plt.show()
```



Decision Tree Classifier

```
In [29]: # Importing necessary packages
from sklearn.tree import DecisionTreeClassifier
```

```
In [30]: # Initializing the model
dtc = DecisionTreeClassifier(random_state=30)
```

```
In [31]: # Training the model
dtc.fit(x_train, y_train)
```

Out[31]: DecisionTreeClassifier(random_state=30)

In a Jupyter environment, please rerun this cell to show the HTML representation or trust the notebook.

On GitHub, the HTML representation is unable to render, please try loading this page with nbviewer.org.

```
In [32]: # Testing the model
y_pred1 = dtc.predict(x_test)
```

```
In [33]: # Evaluation metrics
# Accuracy score
accuracy = accuracy_score(y_test, y_pred1)
print("Accuracy of Decision tree model: ", accuracy)
```

Accuracy of Decision tree model: 0.7517006802721088

```
In [34]: # Precision score
precision_yes = precision_score(y_test, y_pred1, pos_label=1)
print("Precision (Yes): " , str(round(precision_yes,2)))
precision_no = precision_score(y_test, y_pred1, pos_label=0)
print("Precision (No): " + str(round(precision_no, 2)))
```

Precision (Yes): 0.27

Precision (No): 0.86

```
In [35]: # Recall score
recall_yes = recall_score(y_test, y_pred1, pos_label=1)
print("Recall (Yes): " + str(round(recall_yes, 2)))
recall_no = recall_score(y_test, y_pred1, pos_label=0)
print("Recall (No): " + str(round(recall_no, 2)))
```

Recall (Yes): 0.29

Recall (No): 0.84

```
In [36]: # F1 score
f1_score_yes = f1_score(y_test, y_pred1, pos_label=1)
print("F1 Score (Yes): " + str(round(f1_score_yes, 2)))
f1_score_no = f1_score(y_test, y_pred1, pos_label=0)
print("F1 Score (No): " + str(round(f1_score_no, 2)))
```

F1 Score (Yes): 0.28

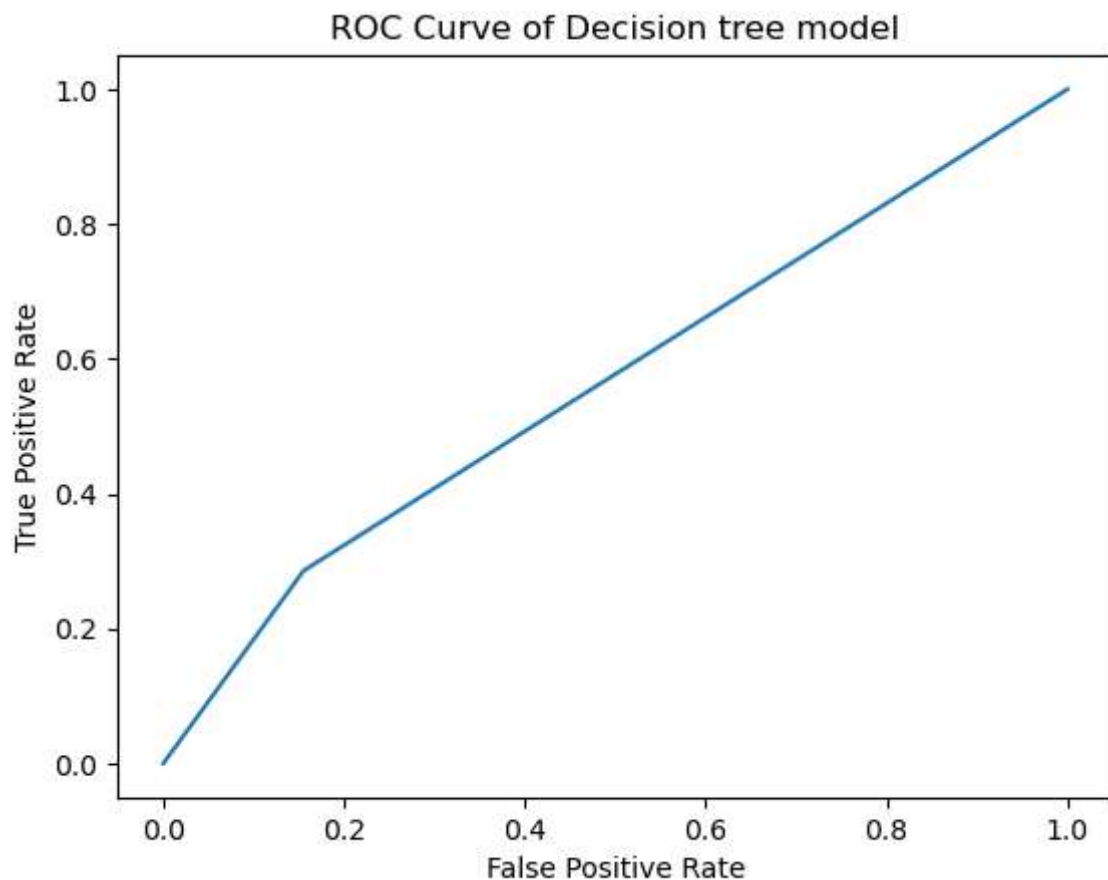
F1 Score (No): 0.85

```
In [37]: # Classification report
print("Classification report of Decision tree model:\n\n",classification_report
```

Classification report of Decision tree model:

	precision	recall	f1-score	support
0	0.86	0.84	0.85	245
1	0.27	0.29	0.28	49
accuracy			0.75	294
macro avg	0.56	0.57	0.56	294
weighted avg	0.76	0.75	0.75	294

```
In [38]: # ROC curve
probability = dtc.predict_proba(x_test)[:,-1]
fpr,tpr,thresholds = roc_curve(y_test,probability)
plt.plot(fpr,tpr)
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
plt.title('ROC Curve of Decision tree model')
plt.show()
```



Random Forest Classifier

```
In [39]: # Importing necessary packages
from sklearn.ensemble import RandomForestClassifier
from sklearn.metrics import accuracy_score
```

```
In [40]: # Initializing the model
rf = RandomForestClassifier(n_estimators=10, criterion='entropy', random_state=30)
```

```
In [41]: # Training the model
rf.fit(x_train, y_train)
```

C:\Users\Admin\AppData\Local\Temp\ipykernel_39480\391630832.py:2: DataConversionWarning: A column-vector y was passed when a 1d array was expected. Please change the shape of y to (n_samples,), for example using ravel().

```
rf.fit(x_train, y_train)
```

Out[41]: RandomForestClassifier(criterion='entropy', n_estimators=10, random_state=30)

In a Jupyter environment, please rerun this cell to show the HTML representation or trust the notebook.

On GitHub, the HTML representation is unable to render, please try loading this page with nbviewer.org.

```
In [42]: rf.score(x_train, y_train)
```

Out[42]: 0.983843537414966

```
In [43]: # Testing the model
y_pred2 = rf.predict(x_test)
```

```
In [44]: # Evaluation metrics
# Accuracy score
accuracy = accuracy_score(y_test, y_pred2)
print("Accuracy of Random forest model: ", accuracy)
```

Accuracy of Random forest model: 0.8435374149659864

```
In [45]: # Precision score
precision_yes = precision_score(y_test, y_pred2, pos_label=1)
print("Precision (Yes): ", str(round(precision_yes, 2)))
precision_no = precision_score(y_test, y_pred2, pos_label=0)
print("Precision (No): " + str(round(precision_no, 2)))
```

Precision (Yes): 0.71
Precision (No): 0.85

```
In [46]: # Recall score
recall_yes = recall_score(y_test, y_pred2, pos_label=1)
print("Recall (Yes): " + str(round(recall_yes, 2)))
recall_no = recall_score(y_test, y_pred2, pos_label=0)
print("Recall (No): " + str(round(recall_no, 2)))
```

Recall (Yes): 0.1
Recall (No): 0.99

```
In [47]: # F1 score
f1_score_yes = f1_score(y_test, y_pred2, pos_label=1)
print("F1 Score (Yes): " + str(round(f1_score_yes, 2)))
f1_score_no = f1_score(y_test, y_pred2, pos_label=0)
print("F1 Score (No): " + str(round(f1_score_no, 2)))
```

F1 Score (Yes): 0.18
F1 Score (No): 0.91

```
In [48]: # Classification Report
print("Classification report of Random Forest model:\n\n", classification_report)
```

Classification report of Random Forest model:

	precision	recall	f1-score	support
0	0.85	0.99	0.91	245
1	0.71	0.10	0.18	49
accuracy			0.84	294
macro avg	0.78	0.55	0.55	294
weighted avg	0.82	0.84	0.79	294


```
In [49]: # ROC curve
probability = rf.predict_proba(x_test)[:,-1]
fpr, tpr, thresholds = roc_curve(y_test, probability)
plt.plot(fpr, tpr)
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
plt.title('ROC Curve of Random forest model')
plt.show()
```

